

Artificial Intelligence

Data Mining, classification

Paweł Kasprowski, PhD, DSc.



Artificial Intelligence

- When there is no algorithmic solution of the problem
- The application ***learns*** to solve the problem
 - typically using some examples
- Intelligence is the ability to
 - learn based on given data
 - understand
 - use acquired knowledge to solve problems

Problems to be solved

- Classification
 - assign samples to predefined classes
 - [photo of an animal -> species]
 - [text -> author]
 - [scan-path -> novice/expert]
- Regression
 - calculate a value for each sample
 - [images of the house -> price of the house]
 - [eye image -> gaze coordinates]
- Conversion
 - convert object into another object
 - [image -> text description of the image]
 - [voice recording -> text of the speech]

Task and method

- Task: find a function $Y = f(X)$ where
 - X – **sample** (input)
 - Y – result – class, value (output) – **label**
- Method (learning by example):
 - take some number of samples X with known output Y (**examples**, labeled samples)
 - build the function based on these examples (learning process)
 - use the function to predict the label for unknown samples

Problem of "generality"

- It is relatively easy to create a function that correctly classifies all examples
- But is this function "general" enough?
 - is it able to correctly classify unknown samples?
- The answer is not trivial and that is why we have a lot of different classification methods
- "No free lunch" theorem

Over-fitting

- If the function (model) is optimized for the given data (given examples) it may have poor generality
- This problem is called over-fitting
- Solution is to simplify the model, e.g. stop learning even when it is not perfect for examples

The simplest algorithms

- K-Nearest Neighbors (kNN)
 - a new sample is classified as majority of its K nearest neighboring examples
- Linear Discriminant Analysis (LDA)
 - finds a linear equation separating positive and negative samples
- Naive Bayes
 - the class of a new sample is calculated based on examples using Bayes equation
- Decision Tree
 - the tree is built using examples by searching for the most discriminative features
 - new samples are going through the tree reaching a leaf with the predicted class

More sophisticated algorithms

- Random Forest
 - a set of randomly generated decision trees voting for the final results
- Support Vector Machines (SVM)
 - recalculates examples into another space where they are easier to separate lineary (kernel trick)
- Neural Networks
 - based on neurons organized into layers

Hyperparameters


- Classification algorithms have parameters
 - (so called hyperparameters)
- kNN – number of neighbors
- Random Forest – number of trees
- SVM – kernel function, parameters of the function, learning rate etc.
- Good choice of hyper parameters influences the quality of the model!
- Bad news: deep learning methods have A LOT hyper parameters

Scikit-learn

- Package implementing:
 - many popular classification algorithms
 - many classic methods for data handling
- General use:
 - load the dataset consisting of samples and their labels
 - create a model using one of the possible algorithms (classes)
 - train the model: `model.fit(samples, labels)`
 - use the model to predict classes:
 - `predicted = model.predict(sample)`

Universal example

python classifier



Creating a model

```
def classification(model_name,samples,labels,rangex,rangey):  
models = {  
    "KNN": KNeighborsClassifier(),  
    "LDA": LinearDiscriminantAnalysis(),  
    "NB": GaussianNB(),  
    "TREE":DecisionTreeClassifier(),  
    "RF":RandomForestClassifier(n_estimators=20),  
    "SVM":SVC(gamma='scale'),  
    "PERC":Perceptron(max_iter=2000),  
    "GB":GradientBoostingClassifier()  
}  
model = models.get(model_name)  
...
```


Using the model

training the model

```
samples = np.array(samples)
```

```
labels = np.array(labels)
```

```
model.fit(samples, labels)
```

classify all points in the matrix

```
all_points = [(x,y) for x in range(rangex) for y in range(rangey)]
```

```
all_points = np.array(all_points)
```

```
r = model.predict(all_points)
```

fill the matrix with predictions

```
result = np.zeros([rangex,rangey])
```

```
for i in range(len(all_points)):
```

```
    result[all_points[i,0],all_points[i,1]]=r[i]
```

```
return result
```

Example for the iris dataset

iris.ipynb

```
import pandas as pd
from sklearn.neighbors import KNeighborsClassifier
data = pd.read_csv('iris.data')
```

```
model = KNeighborsClassifier()
samples = data.values[:,0:4] # first four columns
labels = data.values[:,4] # the fifth column
model.fit(samples, labels)
predicted = model.predict([(5.9, 3.0, 5.1, 1.8)])
print("predicted",predicted)
> predicted ['Iris-virginica']
```

	sl	sw	pl	pw	iris
0	5.1	3.5	1.4	0.2	Iris-setosa
1	4.9	3.0	1.4	0.2	Iris-setosa
2	4.7	3.2	1.3	0.2	Iris-setosa
3	4.6	3.1	1.5	0.2	Iris-setosa
4	5.0	3.6	1.4	0.2	Iris-setosa

Evaluating results

- Hyperparameters may be tuned
 - it is important to be able to check if the model is general
- First rule: never check the model using the same data that you used for training (examples)
 - over-fitting!
- The dataset should always be divided into:
 - training set (used to build the model)
 - test set (to check how it works for unknown samples)

Evaluation

- Dividing into training set and test set may be not enough!
 - we optimize the model (by tuning hyper parameters) for the given test set
- The safest way:
 - training set (for learning)
 - validation set (for hyper parameters tuning)
 - test set (for final evaluation)
 - the "holdout" dataset

Sklearn train_test_split

- `(trainSamples, testSamples, trainLabels, testLabels) = sklearn.model_selection.train_test_split(samples, labels)`
- Two pairs:
 - trainSamples, trainLabels
 - testSamples, testLabels
- Default: 75% train, 25% test
- Parameters:
 - train_size – percent or number of samples
 - test_size – percent or number of samples

iris.ipynb / Train Test Split

Stratification

- Problem: proportion of samples for classes may differ in training and test sets
- For instance: 50-50-50 in the original set
 - training: 40-30-25
 - test: 10-20-25
- Effect: class0 will be treated as more probable!
- Solution: stratified division (maintain distribution)
- Parameter: stratify=labels

Check test samples

iris.ipynb / Check an error

- Checking the test set

```
correct = 0;
```

```
predictedLabels = model.predict(testSamples)
```

```
for i in range(len(testSamples)):
```

```
    print(testLabels[i], "->", predictedLabels[i], end=' ')
```

```
    if(testLabels[i]==predictedLabels[i]):
```

```
        correct = correct + 1; print('OK')
```

```
    else:
```

```
        print('error!!!')
```

```
print("Correct: {} of {} accuracy = {:.2f}"
```

```
    .format(correct, len(testSamples), correct/len(testSamples)))
```

Cross-validation

- Dynamic division into training and test
 - the same examples are sometimes training and sometimes test samples
- 10-fold cross validation:
 - divide the whole dataset into 10 subsets
 - for each subset
 - train the model using the remaining 9 subsets
 - test the results using the chosen subset
 - average the results
- Folds are randomized but may be stratified
 - the same distribution of classes in each fold

Sklearn cross_validate

- ***iris.ipynb / Cross validation***
 - `sklearn.model_selection.cross_validate(model, samples, labels, cv=<number of folds>)`
- Returns for each fold:
 - `fit_time`
 - `score_time`
 - `test_score`
- Many parameters to check other metrics!

Measures

- Accuracy – the number of correctly classified samples to all samples
 - a problem with unbalanced sets
 - if 90% of test samples is positive the blind classifier achieves 90% accuracy
- Precision – the cost of false positives
 - not belonging to class predicted as belonging
- Recall – the cost of false negatives
 - belonging to class predicted as not belonging
- F1-Score – combination of precision and recall
- Cohen's kappa – compares prediction with random prediction

Measures interpretation

- Example: We have built the model that diagnoses COVID-19 based on a genetic test
- For every eye movement sample the model returns:
 - ***H*** - if a person is healthy
 - ***S*** - when a person is sick
- We test the model using
 - 18 samples from sick people (***S***)
 - 82 samples from healthy people (***H***)
- We achieve 82% accuracy
 - is it bad or good?
 - it depends: we must examine a confusion matrix

Confusion matrix

Predicted as >> Real class	H	S
H	81	1
S	17	1

P	N
TP	FN
FP	TN

Accuracy 0.82

H:
precision: 0.83
recall: 0.99
F1-score: 0.90

S:
precision: 0.50
recall: 0.06
F1-score: 0.1

Cohen's kappa: 0.07

Measures

Predicted as >> Real class	H	S
H	80	2
S	15	3

P	N
TP	FN
FP	TN

Accuracy 0.83

H:

precision: 0.84

recall: 0.98

F1-score: 0.90

S:

precision: 0.60

recall: 0.17

F1-score: 0.26

Cohen's kappa: 0.20

Measures

Predicted as >> Real class	H	S
H	66	16
S	2	16

P	N
TP	FN
FP	TN

Accuracy 0.82

H:
precision: 0.97
recall: 0.80
F1-score: 0.88

S:
precision: 0.50
recall: 0.89
F1-score: 0.64

Cohen's kappa: 0.53

Measures

Predicted as >> Real class	H	S
H	64	18
S	0	18

P	N
TP	FN
FP	TN

Accuracy 0.82

H:
precision: 1.00
recall: 0.78
F1-score: 0.88

S:
precision: 0.50
recall: 1.00
F1-score: 0.67

Cohen's kappa: 0.56

Example for three classes

Confusion matrix

[[24 3 9]
 [3 24 24]
 [2 15 49]]

class	precision	recall	f1-score	support
0	0.83	0.67	0.74	36
1	0.57	0.47	0.52	51
2	0.60	0.74	0.66	66
micro avg	0.63	0.63	0.63	153
macro avg	0.67	0.63	0.64	153
weighted avg	0.64	0.63	0.63	153

Accuracy: 0.63

Scikit learn measures

- Useful functions from *sklearn.metrics*:
 - `confusion_matrix(labels, results)`
 - `classification_report(labels, results)`
 - `accuracy_score(labels, results)`
 - `precision_score(labels, results)`
 - `recall_score(labels, results)`
 - `cohen_kappa_score(labels, results)`
 - ...

Measures calculation

iris.ipynb / Measures calculation

```
from sklearn.metrics import classification_report, confusion_matrix,  
    accuracy_score,cohen_kappa_score
```

```
predictedLabels = model.predict(testSamples)  
print(confusion_matrix(testLabels, predictedLabels))  
print(classification_report(testLabels, predictedLabels))  
accuracy = accuracy_score(testLabels, predictedLabels)  
print("Accuracy: {:.2f}".format(accuracy))  
c_kappa = cohen_kappa_score(testLabels, predictedLabels)  
print("Cohen's Kappa: {:.2f}".format(c_kappa))
```

Adding own scorers to cross validation

iris.ipynb / Adding own scorers to cross validation

```
from sklearn.metrics import make_scorer, precision_score, recall_score
```

```
p_scorer = make_scorer(precision_score, average='micro')
```

```
r_scorer = make_scorer(recall_score, average='micro')
```

```
my_scorer = {'precision': p_scorer, 'recall': r_scorer}
```

```
sklearn.model_selection.cross_validate(model, samples, labels, cv=5,  
    scoring=my_scorer)
```

Summary

- Accuracy is often not enough
 - especially for unbalanced datasets
- Cohen's kappa is the reliable way to evaluate results
 - but has some drawbacks as well!
- The evaluation depends on the target we want to achieve
 - False Acceptance Rate ($FAR=0$)
 - no incorrectly classified healthy people
 - False Rejection Rate ($FRR=0$)
 - no incorrectly classified sick people