

Artificial Intelligence

Convolutional Neural Networks

Paweł Kasprowski, PhD, DSc.



Curse of dimensionality

- If the number of features is too big we have the 'curse of dimensionality' problem
 - number of possible 'states' is too big to find any similarities in data
- Typical dataset:
 - A lot of irreleevant features
 - Correlations between features
- For example:
 - Images and pixel values as features
 - 250,000 features for 500x500 images!

Feature extraction

- All previously mentioned algorithms treat input as a set of independent features
 - preferably not correlated
- So the first and most important part of classification is **feature extraction** from real data
- For instance in image classification it could be [1]:
 - **Color** - Color Channel Statistics (Mean, Standard Deviation) and Color Histogram
 - **Shape** - Hu Moments, Zernike Moments
 - **Texture** - Haralick Texture, Local Binary Patterns (LBP)
 - **Others** - Histogram of Oriented Gradients (HOG), Threshold Adjacency Statistics (TAS)

[1] <https://gogul09.github.io/software/image-classification-python>

Convolutional Neural Networks

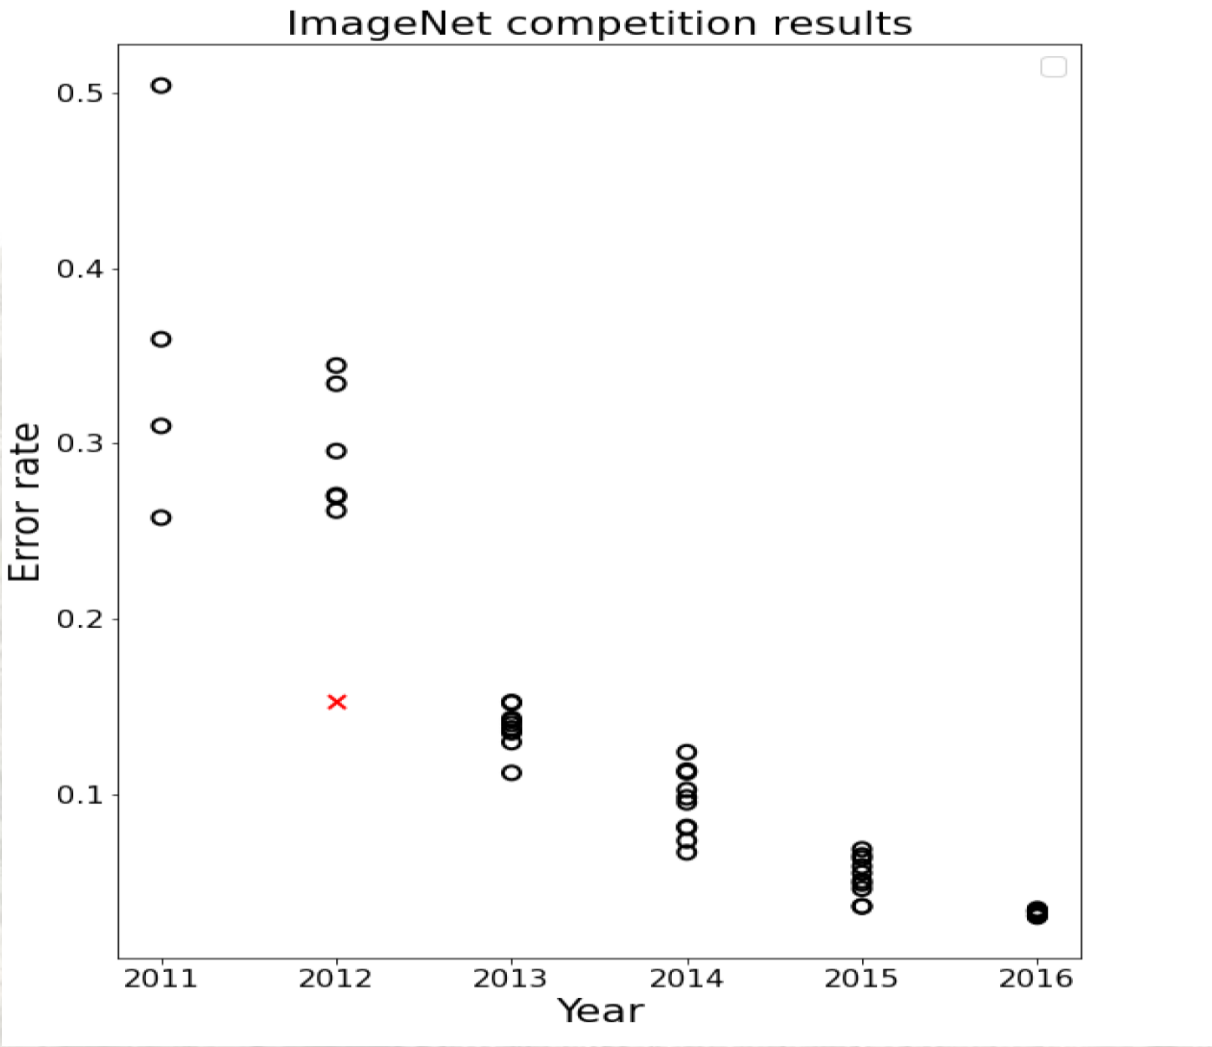
- Network that is not "dense"
 - neuron from layer $N+1$ is not connected to every neuron in layer N
- It preserves "local connectivity"
 - Features (pixels) close to each other are processed together
- Such a network is in fact doing automatic feature extraction in input layers
- Two key properties:
 - not all neurons are connected
 - there are common weights for many connections

Brief history

- 1982 – Kuniyoshi Fukushima, Neocognitron
 - pattern recognition
- 1989 – Yann LeCun, LeNet-5
- 2010 – ImageNet Competition
 - 1000 classes, over million of images
- 2012 – Alex Krizhevsky, George Hinton: AlexNet
 - 15% error rate for ImageNet (runner-up: 26%)
 - 8 layers
- 2015 – Deep Residual Nets wins ImageNet
 - over 150 layers!

AlexNet

- Krizhevsky, Alex, Ilya Sutskever, and Geoffrey E. Hinton. "Imagenet classification with deep convolutional neural networks." *Communications of the ACM* 60.6 (2017): 84-90.
 - [Cited by 73,253](#)



<https://m.marefa.org/>

Why now?

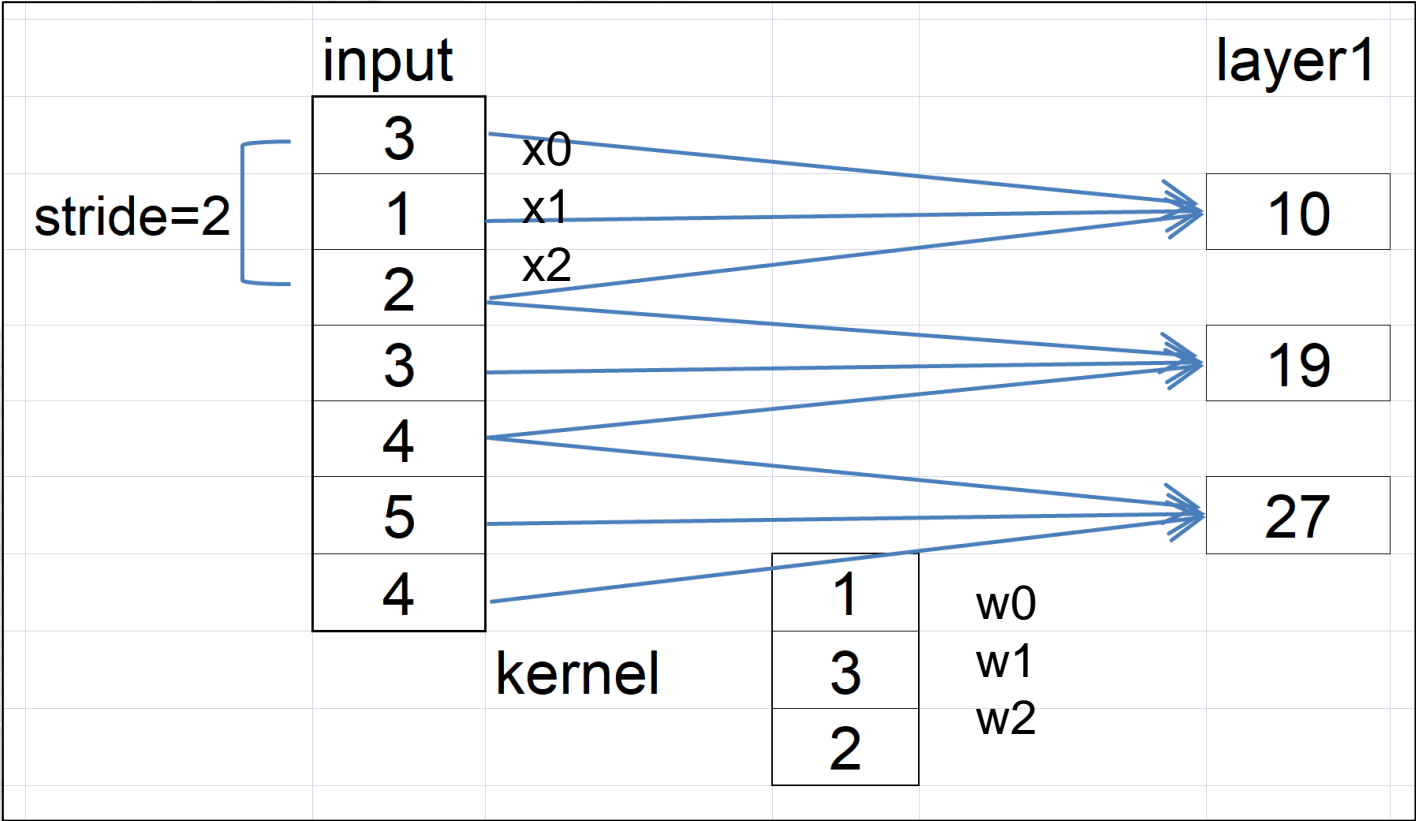
- Neural Network training and predicting involves a lot of calculations
 - Only new computers are capable to calculate it in the reasonable time (esp. with CUDA/GPU)
- Neural Network training requires a lot of training data
 - Huge datasets like ImageNet were not available before
 - Everybody may easily create their own dataset using internet resources
- Some advancements in algorithms
 - optimizers, learning rate adjustments, RELU,...

Convolutional Filter

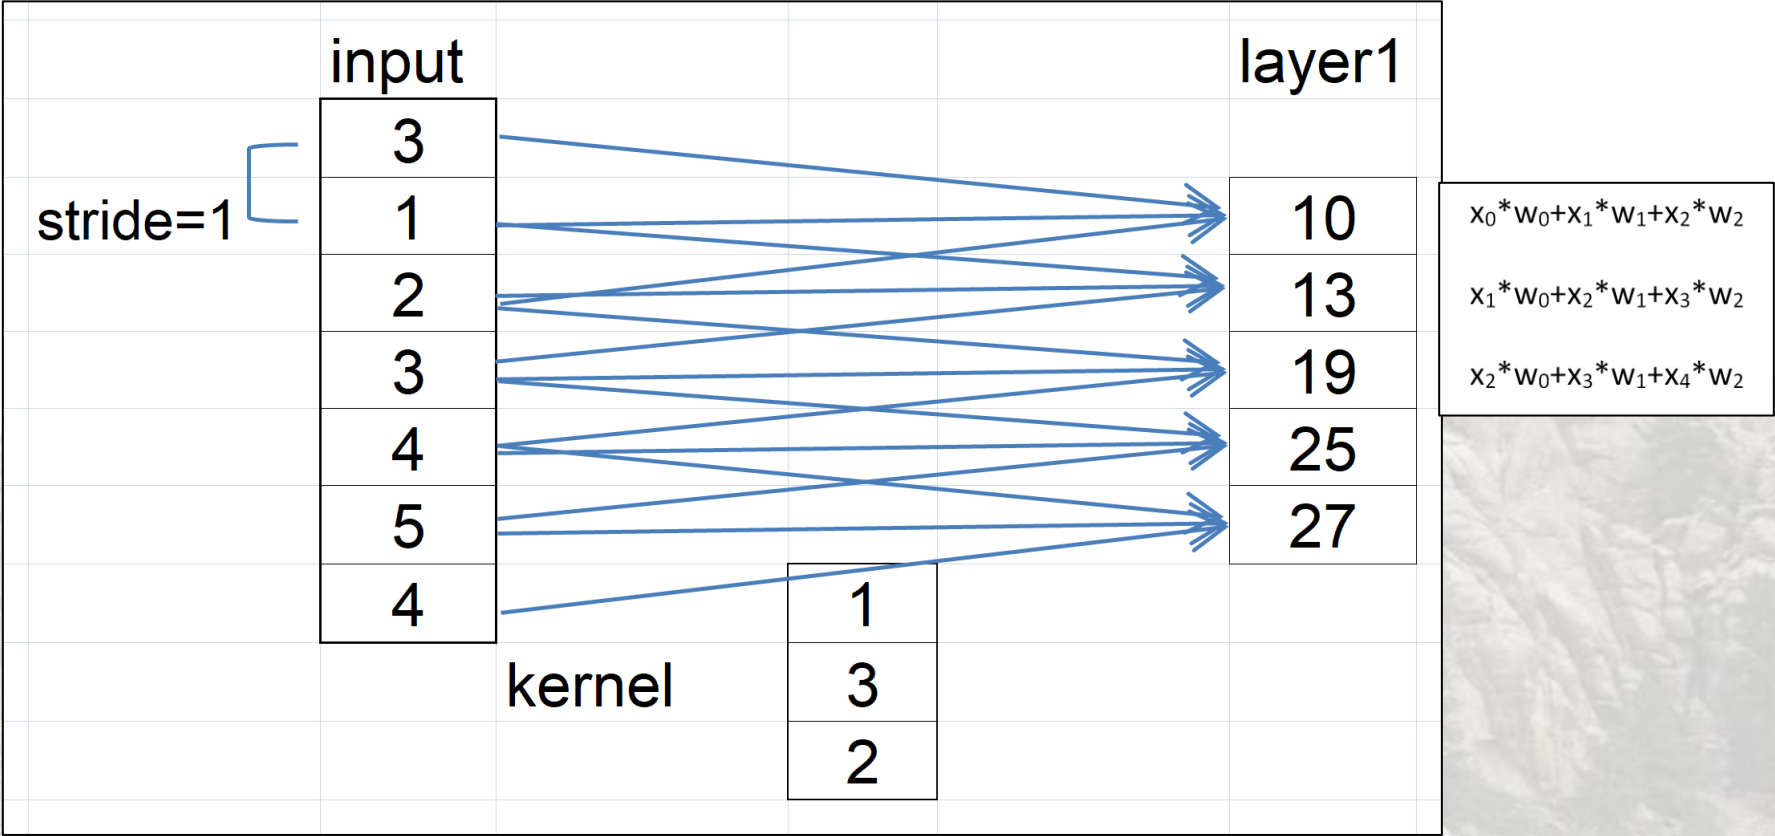
- The filter (matrix $N \times M$) applied to each pixel and its neighborhood
 - sharpening filters
 - smoothing filters
 - edge detection filters
- Examples:
 - <https://setosa.io/ev/image-kernels/>

CNN - 1D example

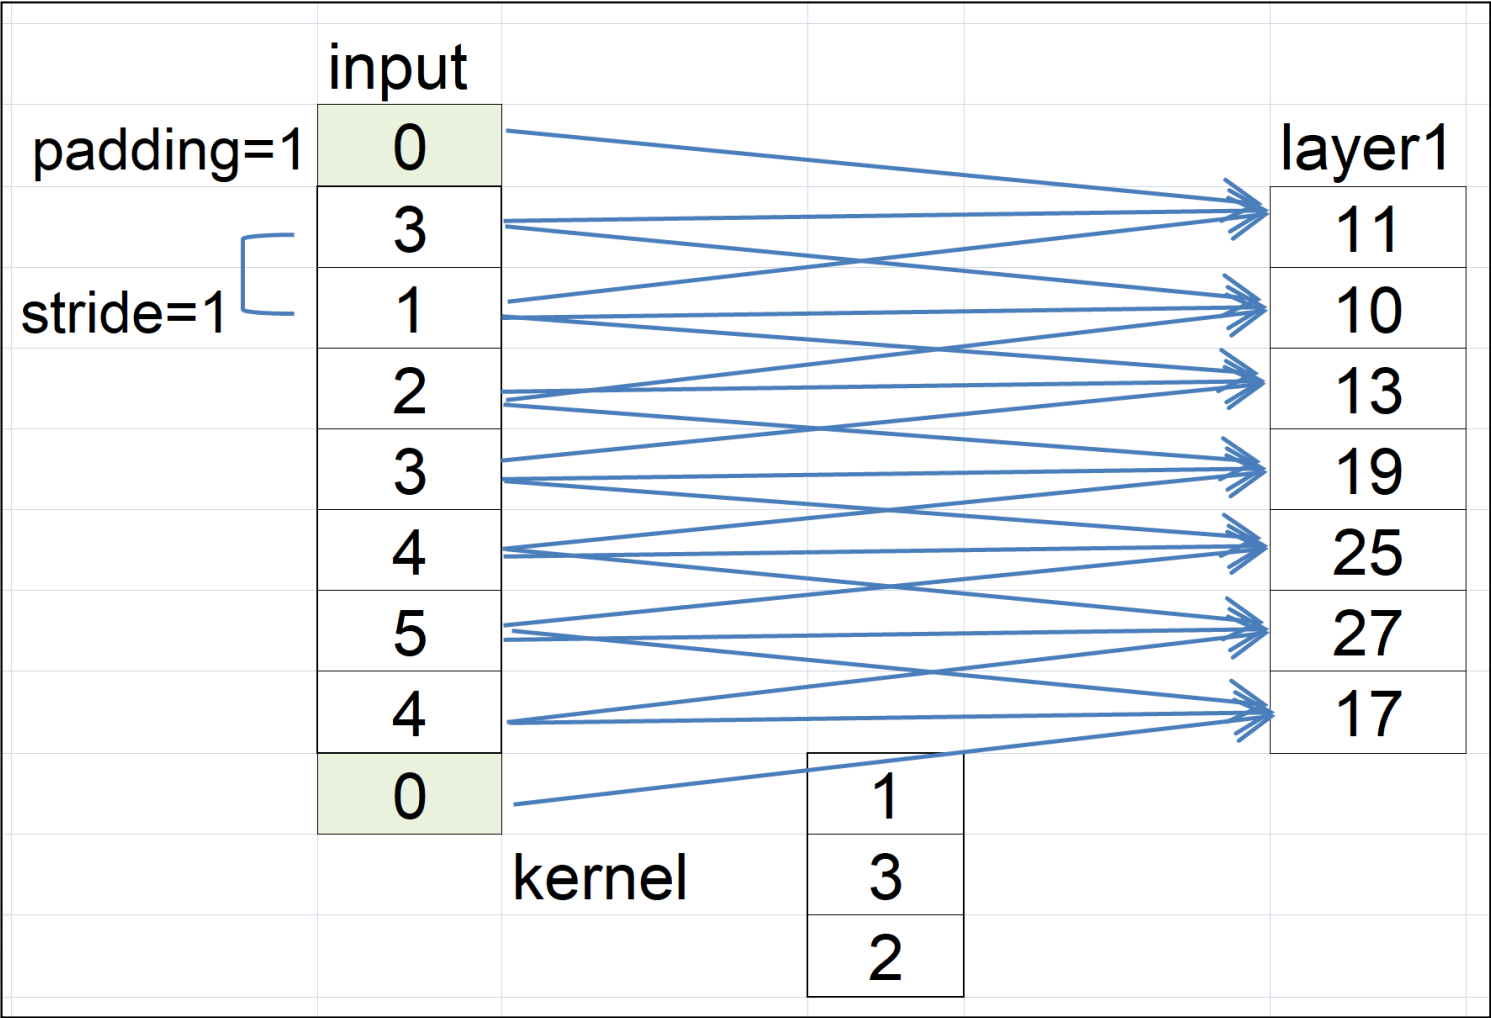
$$x_0 * w_0 + x_1 * w_1 + x_2 * w_2$$



CNN - stride=1

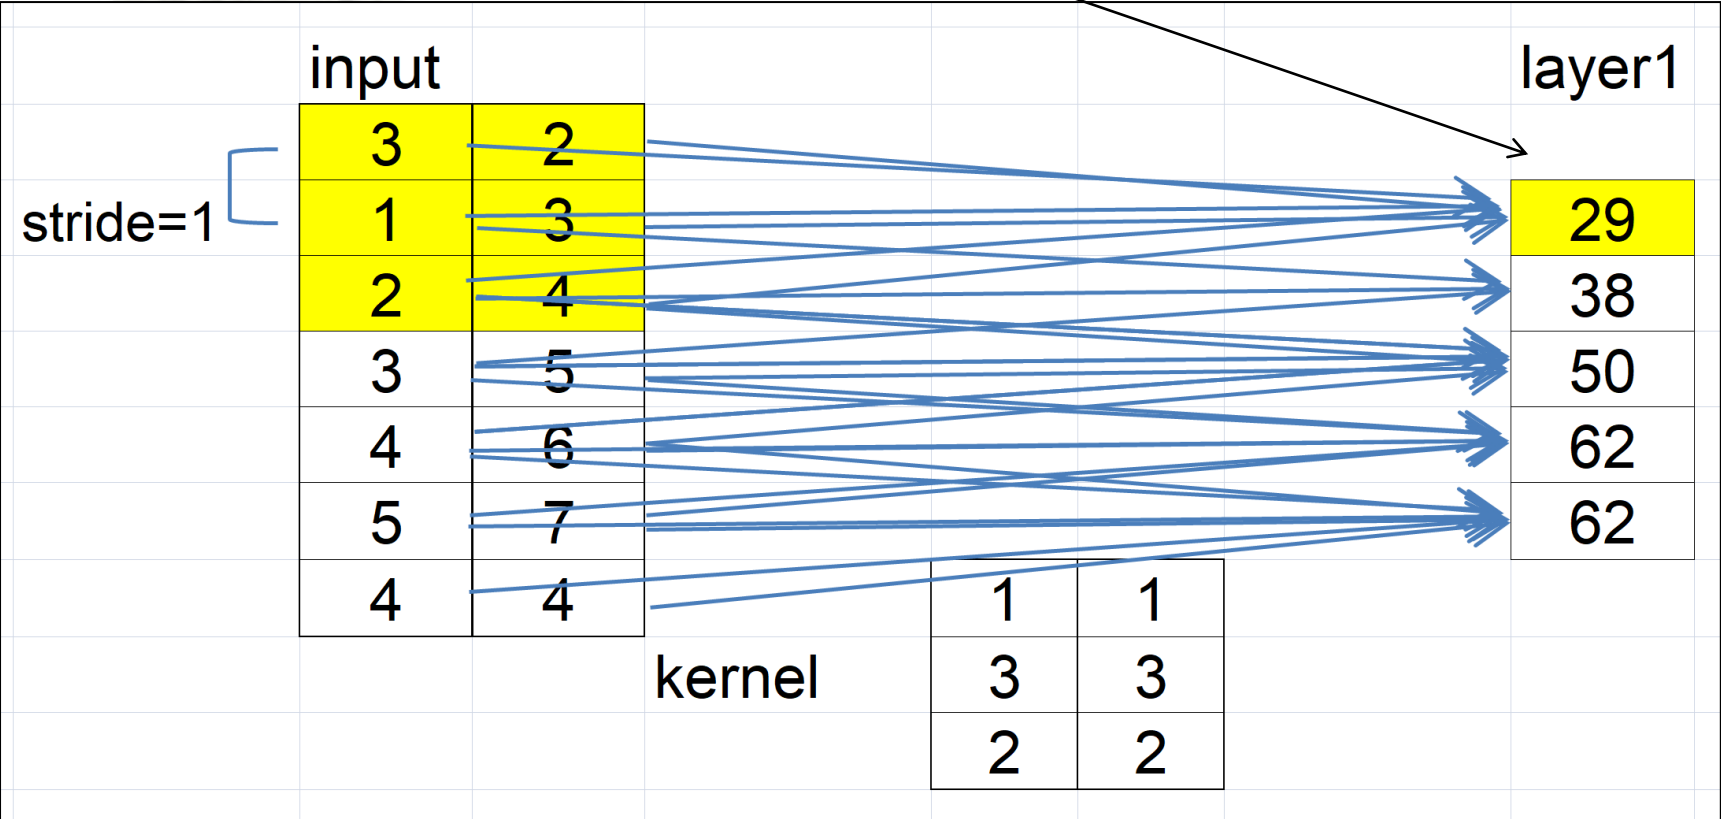


CNN - stride=1, padding=1

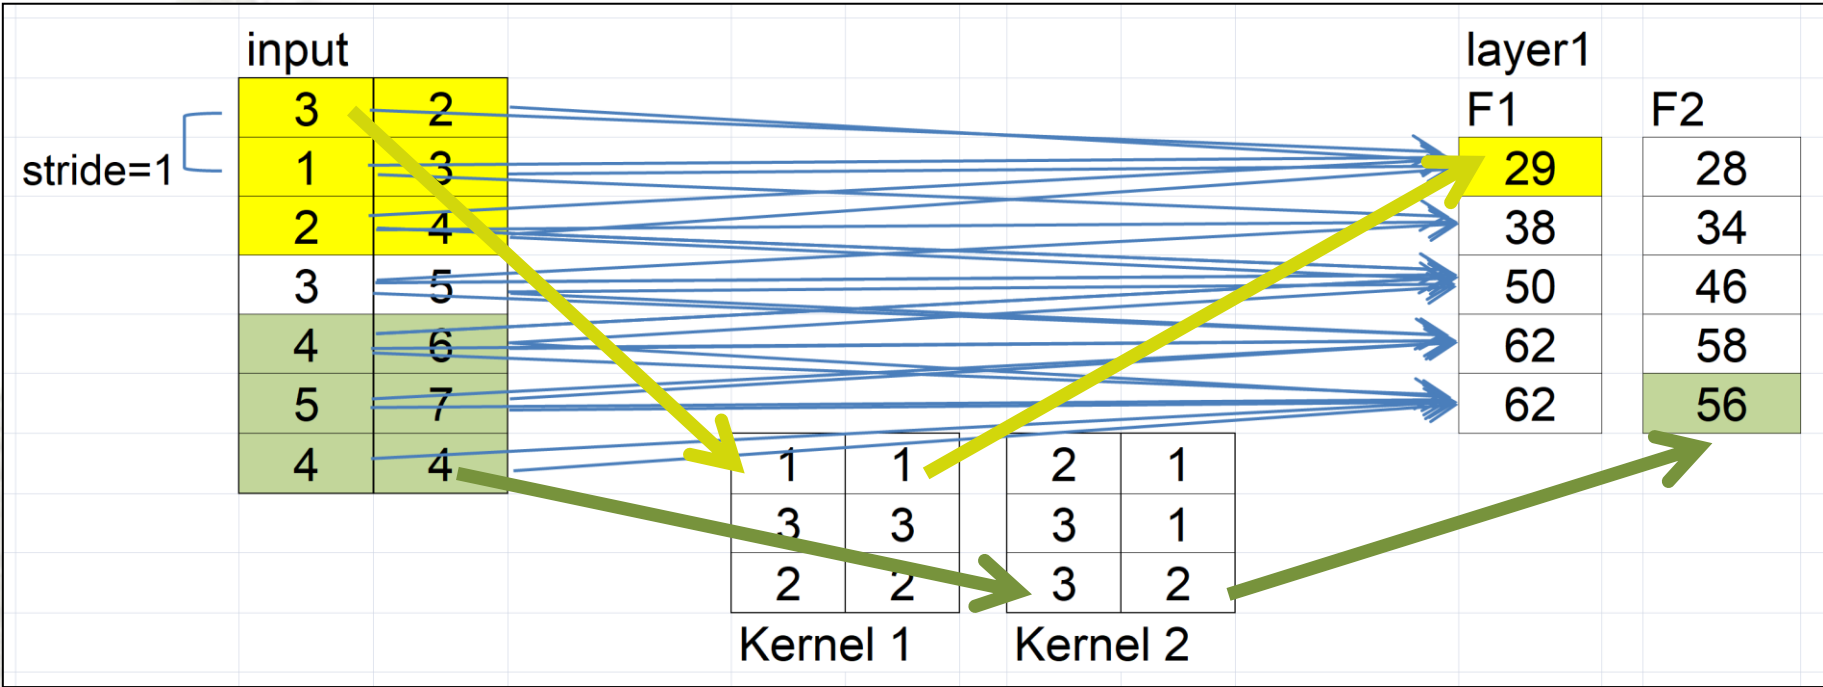


CNN - 2D

$$X_{00} * W_{00} + X_{01} * W_{01} + X_{02} * W_{02} + X_{10} * W_{10} + X_{11} * W_{11} + X_{12} * W_{12}$$



CNN – 2D, many filters



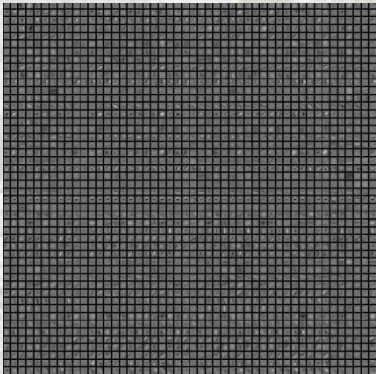
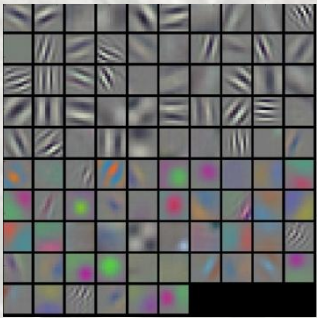
Convolutional Layer

- Parameters:
 - filters – how many kernels
 - kernel size – 1D or 2D
 - stride – step for applying the kernel
 - padding – add borders with zeroes
 - VALID – no padding
 - SAME – padding to preserve size (if STRIDE=1)
- Keras:
 - `Conv2D(filters=64, kernel_size=(3,3), padding=SAME, input_shape=(256,256))`

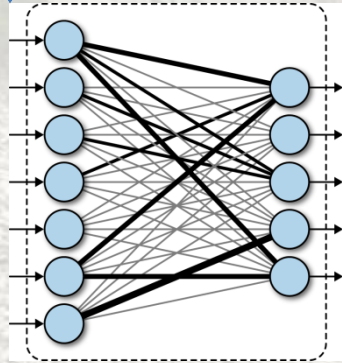
Advantages of CNN

- Takes into account spatial/temporal relationships between features
- May find patterns in data
- Builds own filters that extract useful information
- The output from convolutional layer is a set of filters representing various properties of the image
 - i.e. features! – which are automatically created

Example of CNN



Feature extraction
finished here



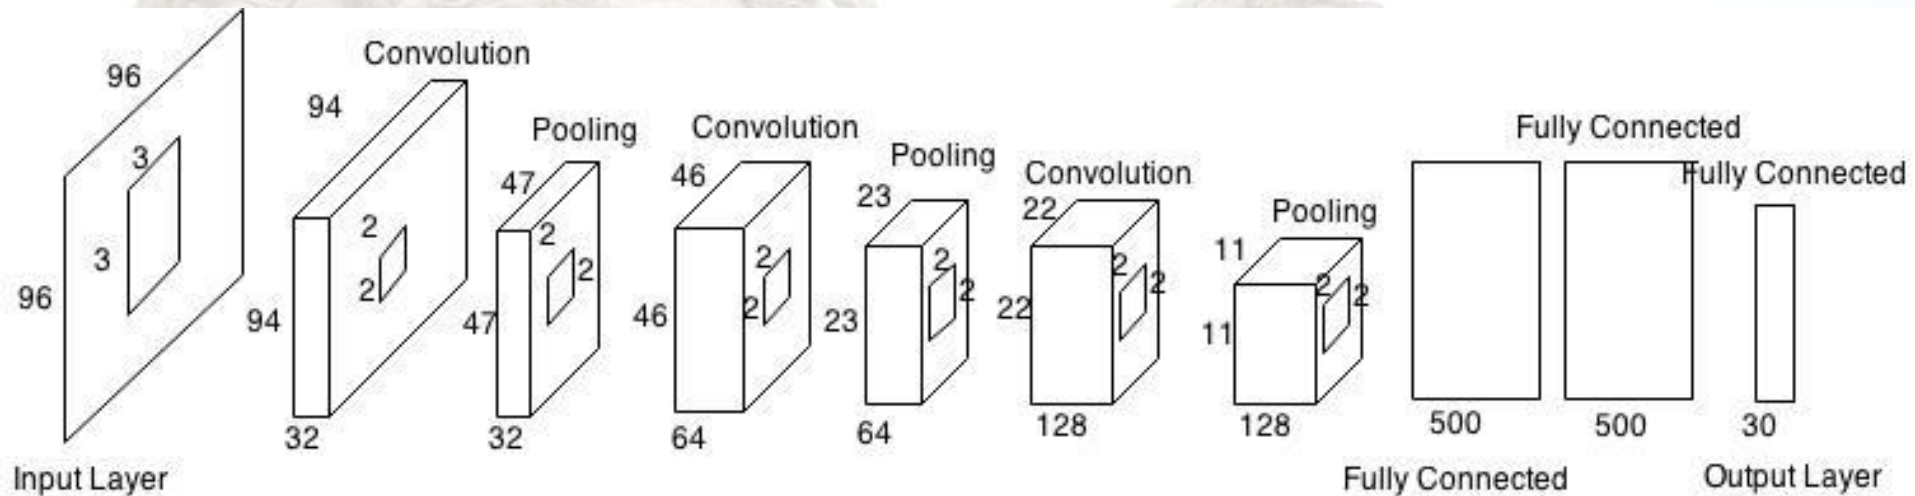
cat:	0.8
dog:	0.2
house:	0.01
car:	0.05
swan:	0.04

CNN Layers in Keras

- `Conv2D(filters=16, kernel_size=(3, 3), padding="same", input_shape=(120,120))`
 - classic layer
- `MaxPooling2D(pool_size=(2, 2))`
 - calculate max for given area
 - reduces size
- `Dropout(rate=0.25)`
 - randomly set rate percent of weights to zero
 - helps to prevent overfitting
- `BatchNormalization(axes=-1)`
 - normalizes output from the layer

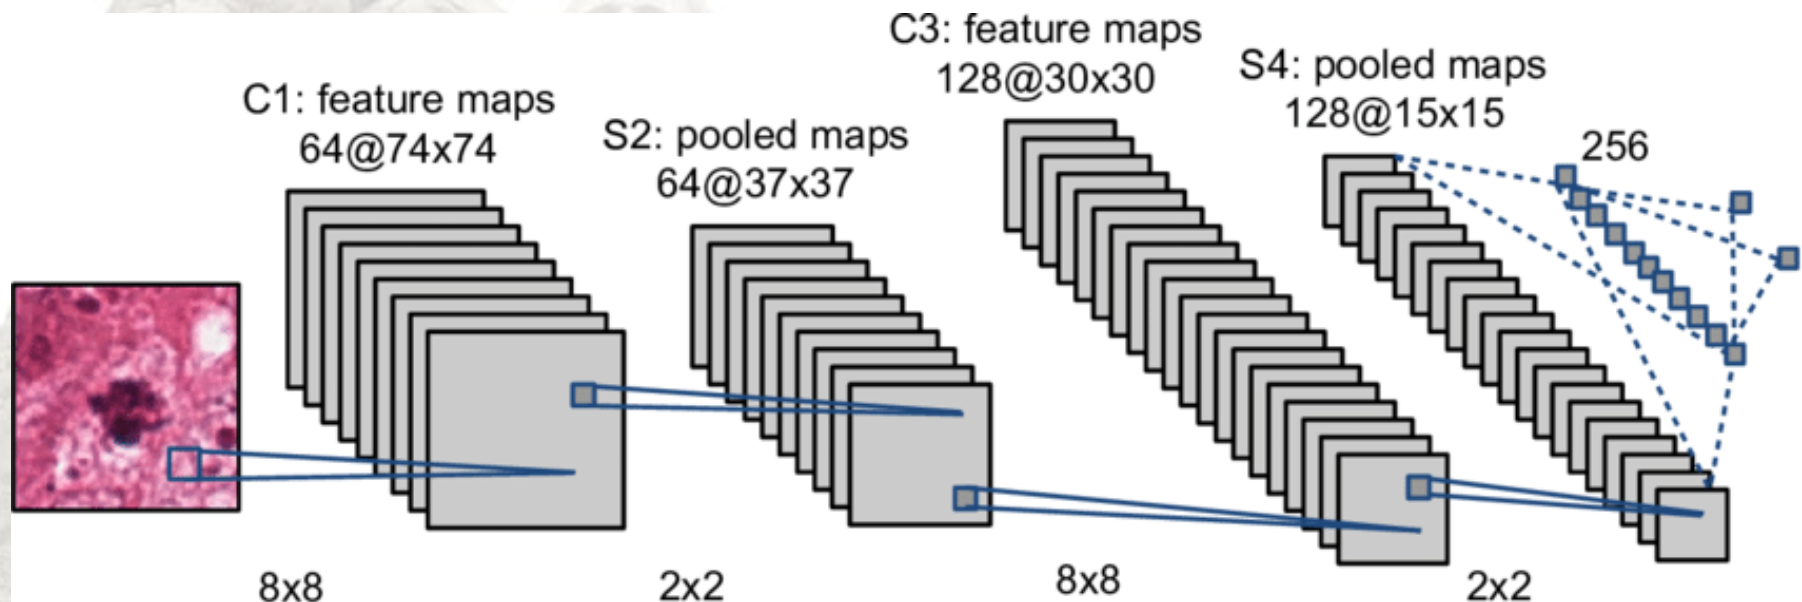
Cascade of layers

- Conv2D>MaxPooling>Conv2D>MaxPooling>Dropout
- Example:



<https://www.hackerearth.com/practice/notes/>

Another example



<https://www.researchgate.net/publication/266734716>

Flowers as images

flowers.ipynb

- Classification using:
 - Classic ANN
 - Decision Tree
 - Convolutional Neural Network